

SEEM5020 Algorithms for Big Data

Nearest Neighbor Search: Locality Sensitive Hashing

Sibo WANG

Department of Systems Engineering and Engineering Management
The Chinese University of Hong Kong

Nearest Neighbor Search (NNS)

Definition 1 (Nearest Neighbor Search)

Given a set $P = \{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n\}$ of n points in $\mathbb{R}^d$ and a distance metric $\text{dist}(\cdot, \cdot)$, for any query point $\mathbf{q} \in \mathbb{R}^d$, the nearest neighbor search query finds the **point closest** to $\mathbf{q}$ in P according to the provided distance metric.

Definition 2 (r -near neighbor search (r -NNS))

Still consider the input set P of n d -dimensional points and a distance metric $\text{dist}(\cdot, \cdot)$ is given. For any query point $\mathbf{q}$, the r -near neighbor search (if exists) returns **a point** $\mathbf{x} \in P$ s.t. $\text{dist}(\mathbf{q}, \mathbf{x}) \leq r$

The r -near neighbor search problem can be treated as a decision problem of the nearest neighbor search problem. To solve the nearest neighbor search, we can apply the decision version with $\log(d_{\max}/d_{\min})$ iterations.

Nearest Neighbor Search: Exact Solutions

Consider the Euclidean distance, i.e., $\text{dist}(\mathbf{x}, \mathbf{y}) = \|\mathbf{x} - \mathbf{y}\|_2$.


- When $d = 1$: sort the data and then for any input query point, we can do a binary search to find the closest point.
 - $O(n)$ space and $O(\log n)$ query time.
- When $d = 2$: Building a Voronoi diagram.
 - $O(n)$ space and $O(\log n)$ query time.
- When $d > 2$:
 - Voronoi diagram: $O(n^{\lceil d/2 \rceil})$ space. Too expensive!
 - Linear search: $O(d \cdot n)$ search time.

Approximate Nearest Neighbor Search

Definition 3 (c -Approximate Nearest Neighbor Search (c -ANNS))

Given a set $P = \{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n\}$ of n points in $\mathbb{R}^d$ and a distance metric $\text{dist}(\cdot, \cdot)$, for any query point $\mathbf{q} \in \mathbb{R}^d$, the c -approximate nearest neighbor search query returns **an arbitrary point** $\mathbf{x}$ so that $\text{dist}(\mathbf{q}, \mathbf{x}) \leq c \cdot \text{dist}(\mathbf{q}, \mathbf{x}^*)$, where $\mathbf{x}^*$ is the nearest neighbor of $\mathbf{q}$.

- Point $\mathbf{x}_4$ is the nearest neighbor of the query point $\mathbf{q}$.
- Within the range of $2 \cdot \text{dist}(\mathbf{q}, \mathbf{x}_4)$, $\mathbf{x}_1, \mathbf{x}_2, \mathbf{x}_3$ all fall into it.
- We can return any of $\mathbf{x}_1, \mathbf{x}_2$, and $\mathbf{x}_3$, which is a valid answer for the 2-ANNS query.



Approximate r -Near Neighbor Search

Definition 4 (c -Approximate r -Near Neighbor Search (c, r) -ANNS)

For any query point $\mathbf{q}$, if there exists a point $\mathbf{x}$ in P such that $\text{dist}(\mathbf{q}, \mathbf{x}) \leq r$, then the c -approximate r -near neighbor search query returns a point $\mathbf{x}' \in P$ so that $\text{dist}(\mathbf{q}, \mathbf{x}') \leq c \cdot \text{dist}(\mathbf{q}, \mathbf{x})$.

Similarly, we can answer the c -approximate nearest neighbor query via a binary search on the radius r with the (c, r) -ANNS query. Next, we focus on this (c, r) -ANNS query, which will be solved by locality Sensitive hashing.

Table of Contents

1 Locality Sensitive Hashing

2 LSH Family for Different Distance Measures


- LSH for Euclidean distance
- LSH for Cosine Distance
- LSH for Jaccard Distance (for Binary Vectors)

Locality-Sensitive Hashing [1]

Intuition: For two points $\mathbf{x}$ and $\mathbf{y}$, locality-sensitive hashing will hash the close data points into the same buckets with a higher probability.

Consider a simple idea that projects the data points into a random line crossing the origin.

- Close points $\mathbf{x}_1$ and $\mathbf{x}_2$ are projected into the same bucket with ID 3.
- The far point $\mathbf{x}_4$ of $\mathbf{x}_1$ is in a different bucket, with ID 1.
- The far point $\mathbf{x}_5$ is in the same bucket as $\mathbf{x}_1$. But the chances of such events are much lower.



Locality-Sensitive Hashing (LSH) Families

Definition 5 ($(r, c \cdot r, p_1, p_2)$ -sensitive)

Given a distance measure $\text{dist}(\cdot, \cdot)$, a fixed r , a family $\mathcal{H}$ of hash functions is said to be $(r, c \cdot r, p_1, p_2)$ -sensitive, where $p_1 > p_2$ and $c > 1$, if h randomly drawn from $\mathcal{H}$ satisfies the following:

- $\mathbb{P}[h(\mathbf{x}) = h(\mathbf{y})] \geq p_1$ when $\text{dist}(\mathbf{x}, \mathbf{y}) \leq r$.
 - If $\mathbf{x}$ and $\mathbf{y}$ are close, the collision probability is high.
- $\mathbb{P}[h(\mathbf{x}) = h(\mathbf{y})] \leq p_2$ when $\text{dist}(\mathbf{x}, \mathbf{y}) \geq c \cdot r$.
 - If $\mathbf{x}$ and $\mathbf{y}$ are sufficiently far, the probability of collision is low.

A key parameter: the gap between p_1 and p_2 measured as $\rho = \frac{\log p_1}{\log p_2}$. We will see the role of ρ later in our analysis.


Here, we assume that we have already had the locality-sensitive family $\mathcal{H}$ for the distance measure dist . We will see how to design the LSH family for different metrics later.

The Issue with a Single LSH Function from $\mathcal{H}$

Given a query point $\mathbf{q}$, even though there is a small probability p_2 that a point $\mathbf{x}$ will collide with $\mathbf{q}$ if the distance is **far**, i.e. at least $c \cdot r$, there might exist $O(n)$ such **far** data points.

- There exist $O(p_2 \cdot n)$ **far** points that collide with $\mathbf{q}$ in expectation.
- These **far** data points that collide with $\mathbf{q}$ are called **false positives**.

- For query point $\mathbf{q}$, x_5 is a false positive



Pruning False Positives with AND Operation

We need to reduce the number of false positives!

- Solution: Pick multiple independent hash functions $h_1, h_2, \dots, h_k$ from $\mathcal{H}$. Define a new hash function $g(\mathbf{x})$ with the AND operation:

$$g(\mathbf{x}) = \langle h_1(\mathbf{x}), h_2(\mathbf{x}), \dots, h_k(\mathbf{x}) \rangle$$

where $g(\mathbf{x}_1) = g(\mathbf{x}_2)$ if and only if

$$h_1(\mathbf{x}_1) = h_1(\mathbf{x}_2) \wedge h_2(\mathbf{x}_1) = h_2(\mathbf{x}_2) \wedge \dots \wedge h_k(\mathbf{x}_1) = h_k(\mathbf{x}_2)$$

Lemma 1

Given a family $\mathcal{H}$ of $(r, c \cdot r, p_1, p_2)$ -sensitive hash functions, by randomly choosing k hash functions from $\mathcal{H}$ and defining it as

$$g(\mathbf{x}) = \langle h_1(\mathbf{x}), h_2(\mathbf{x}), \dots, h_k(\mathbf{x}) \rangle,$$

$g(\mathbf{x})$ constitutes a family $\mathcal{G}$ of $(r, c \cdot r, p_1^k, p_2^k)$ -sensitive hash functions.

Choosing the Appropriate k

We choose k so that the expected number of **far** points is ≤ 1 .

- This can be done by setting $p_2^k = \frac{1}{n} \rightarrow k = \log_{1/p_2}(n)$.
- Accordingly, $p_1^k = 1/n^p$. If there is only one point that is the r -near neighbor of the query point $\mathbf{q}$, then in expectation we only have the probability $\frac{1}{n^p}$ that the point hashes to the same bucket as the query point $\mathbf{q}$.
- As we need to spend at least $O(1)$ time to do a table lookup.
 - The cost of the false positive hence can be bounded by the table lookup cost, without incurring additional cost.

Multiple Hash Tables

We now need to increase the chance that a r -near neighbor of $\mathbf{q}$ hashes to the same bucket as $\mathbf{q}$.

- We can choose a sufficiently large number L of hash functions from $(r, c \cdot r, p_1^k, p_2^k)$ -sensitive hash family $\mathcal{G}$.
- Then, the r -near neighbors will hash to **at least one of these L hash functions** with high probability.
 - We set $L = n^\rho$. The reason will be explained shortly.
 - The storage is then: $O(n \cdot L) = O(n^{1+\rho})$.

Query Processing

The LSH includes L hash tables, where for each hash table, the hash function is chosen from the $(r, c \cdot r, p_1^k, p_2^k)$ -sensitive family $\mathcal{G}$. To answer a query, it proceeds as follows:

- Retrieve data points from buckets $g_1(\mathbf{q}), g_2(\mathbf{q}), \dots, g_L(\mathbf{q})$ one by one.
- For each retrieved data point $\mathbf{x}$, we compute the distance $\text{dist}(\mathbf{q}, \mathbf{x})$.
- Then, (i) we return the first data point such that the distance to $\mathbf{q}$ is no larger than $c \cdot r$, or (ii) we have retrieved all data points from the L buckets but no point within distance $c \cdot r$, then we return **failure**.

What if there are too many data points in these L buckets? How to bound the search complexity?

- **Fix**: we stop the search when we have retrieved $3L$ data points (but no cr -near neighbors) and we return **failure**.
- This bounds the search complexity to $O(d \cdot k \cdot L)$.

Theoretical Analysis

For a query point $\mathbf{q}$, the LSH query algorithm answers **correctly**:

- If there is no point that is a $c \cdot r$ -near neighbor of $\mathbf{q}$, hence the query returns **failure**.
- If the LSH query algorithm returns a data point $\mathbf{x}$, its distance to the query point $\mathbf{q}$ is always bounded by $c \cdot r$.

Theoretical Analysis

For a query point $\mathbf{q}$, the LSH query algorithm answers **correctly**:

- If there is no point that is a $c \cdot r$ -near neighbor of $\mathbf{q}$, hence the query returns **failure**.
- If the LSH query algorithm returns a data point $\mathbf{x}$, its distance to the query point $\mathbf{q}$ is always bounded by $c \cdot r$.

How to bound the failure probability when an r -near neighbor $\mathbf{x}$ of $\mathbf{q}$ exists. We consider the following two **bad events**:

- Event E_1 : For any r -near neighbor $\mathbf{x}$, it does not collide with query point $\mathbf{q}$.
- Event E_2 : There are too many (more than $3L$) **far** points colliding with $\mathbf{q}$ in these L hash functions.

Theoretical Analysis (Cont.)

Lemma 2

Event E_1 occurs with probability no more than $\frac{1}{e}$.

Proof.

We analyze the case where there is only one r -near neighbor $\mathbf{x}$ of the query point $\mathbf{q}$. Obviously, the more r -near neighbors $\mathbf{q}$ has, the smaller the probability of event E_1 will be.

$$\begin{aligned}\mathbb{P}[E_1] &= \mathbb{P}[g_1(\mathbf{x}) \neq g_1(\mathbf{q}) \wedge g_2(\mathbf{x}) \neq g_2(\mathbf{q}) \cdots g_L(\mathbf{x}) \neq g_L(\mathbf{q})] \\ &= \mathbb{P}[g_1(\mathbf{x}) \neq g_1(\mathbf{q})] \cdot \mathbb{P}[g_2(\mathbf{x}) \neq g_2(\mathbf{q})] \cdots \mathbb{P}[g_L(\mathbf{x}) \neq g_L(\mathbf{q})] \\ &\leq \left(1 - \frac{1}{n^\rho}\right) \cdot \left(1 - \frac{1}{n^\rho}\right) \cdots \left(1 - \frac{1}{n^\rho}\right) = \left(1 - \frac{1}{n^\rho}\right)^L \text{ (recap: } L = n^\rho) \\ &= (1 - 1/L)^L \leq \frac{1}{e}\end{aligned}$$

This finishes the proof. □

Theoretical Analysis (Cont.)

Lemma 3

Event E_2 occurs with a probability no more than $\frac{1}{3}$.

Proof.

Let $\mathbf{x}$ be a bad point such that $\text{dist}(\mathbf{q}, \mathbf{x}) > c \cdot r$. Let Y be a random variable to indicate the number of bad points examined. As we have shown, for a given hash function $g_i(\cdot)$, the expected number of bad points that collide with $\mathbf{q}$ is bounded by 1. Thus, the expected number of points that collide with the L hash functions is bounded by L . Then, by Markov's inequality, we have:

$$\mathbb{P}[Y \geq 3L] \leq \frac{\mathbb{E}[Y]}{3L} \leq \frac{L}{3L} = \frac{1}{3}.$$

This finishes the proof. □

Theoretical Analysis (Cont.)

Theorem 1

The LSH query algorithm correctly returns an c -approximate r -near neighbor with probability at least $\frac{2}{3} - \frac{1}{e}$.

The above theorem can be derived via a combination of Lemmas 2-3 and the union bound.

Consider the following questions:

- How to increase the success probability to a larger constant?
- In case we want to return the c -approximate nearest neighbor query with a success probability of at least $\frac{2}{3} - \frac{1}{e}$, how should we set the parameters?
- In the median trick, we need to have the success probability larger than $1/2$ to boost the probability. Here, do we need to have the same constraint? Why?

Table of Contents

1 Locality Sensitive Hashing

2 LSH Family for Different Distance Measures

- LSH for Euclidean distance
- LSH for Cosine Distance
- LSH for Jaccard Distance (for Binary Vectors)

LSH for Euclidean Distance [2]

Intuition: Projection onto random lines (crossing the origin) and divide them into different buckets.


- Version 1: We choose a random line by randomly sampling a unit vector $\mathbf{u}$ and then divide them into different buckets.

$$h_{\mathbf{u}}(\mathbf{x}) = \lceil \frac{\mathbf{u} \cdot \mathbf{x}}{r} \rceil.$$

Issue: Even if two points are very close, we might still project them into different buckets.

- Final version: We add a random offset $b \in [0, r]$ to the result.

$$h_{\mathbf{u},b}(\mathbf{x}) = \lceil \frac{\mathbf{u} \cdot \mathbf{x} + b}{r} \rceil. \quad (1)$$



Generating the Random Gaussian Vector $\mathbf{u}$

Next, we show how to generate a random Gaussian vector in $\mathbb{R}^d$:

- Pick d independent and identically distributed (i.i.d) random variables $Z_1, Z_2, \dots, Z_d$ from Gaussian distribution $\mathcal{N}(0, 1)$. Let $\mathbf{u} = (Z_1, Z_2, \dots, Z_d)$.
- The vector $\mathbf{u}$ is also called a **random Gaussian vector**.

Recap: **Property** of Gaussian distributions.

Lemma 4

Assume that we have two random variables X and Y that are sampled from two independent Gaussian distributions, i.e., $X \sim \mathcal{N}(\mu_1, \sigma_1^2)$ and $Y \sim \mathcal{N}(\mu_2, \sigma_2^2)$. Then, their sum $Z = X + Y$ also follows a Gaussian distribution specified as $\mathcal{N}(\mu_1 + \mu_2, \sigma_1^2 + \sigma_2^2)$.

Analysis of the Hashing Scheme

Recap that for the hash function:

- We pick a bucket width $r > 0$
- a number b sampled from $(0, r)$ uniformly at random
- and a random Gaussian vector $\mathbf{u}$.

The final hash function is $h_{\mathbf{u},b}(\mathbf{x}) = \lceil \frac{\mathbf{x} \cdot \mathbf{u} + b}{r} \rceil$. We omit the subscript and denote it as $h(\mathbf{x})$ directly when the context is clear.

Lemma 5

Given a vector $\mathbf{x}$ and a random Gaussian vector $\mathbf{u}$, then $Z = \mathbf{x} \cdot \mathbf{u}$ follows a Gaussian distribution $\mathcal{N}(0, (\|\mathbf{x}\|_2)^2)$ and $\mathbb{E}[Z^2] = (\|\mathbf{x}\|_2)^2$.

Analysis of the Hashing Scheme (Cont.)

Given the hashing scheme $h(\mathbf{x})$, now we consider the probability:

- $p_1 = \mathbb{P}[h(\mathbf{x}_1) = h(\mathbf{x}_2)]$ if $\text{dist}(\mathbf{x}_1, \mathbf{x}_2) = \|\mathbf{x}_1 - \mathbf{x}_2\|_2 \leq r$.
- $p_2 = \mathbb{P}[h(\mathbf{x}_1) = h(\mathbf{x}_2)]$ if $\text{dist}(\mathbf{x}_1, \mathbf{x}_2) = \|\mathbf{x}_1 - \mathbf{x}_2\|_2 \geq c \cdot r$

Define $\mathbf{x}' = \mathbf{x}_1 - \mathbf{x}_2$.

$$h(\mathbf{x}_1) = h(\mathbf{x}_2) \Leftrightarrow \lceil \frac{\mathbf{x}_1 \cdot \mathbf{u} + b}{r} \rceil = \lceil \frac{\mathbf{x}_2 \cdot \mathbf{u} + b}{r} \rceil$$

As we have analyzed in Lecture 1,

- if $|\mathbf{x}_1 \cdot \mathbf{u} - \mathbf{x}_2 \cdot \mathbf{u}| \leq r$, the probability that $h(\mathbf{x}_1) = h(\mathbf{x}_2)$ holds is $1 - |\mathbf{x}_1 \cdot \mathbf{u} - \mathbf{x}_2 \cdot \mathbf{u}|/r$.
- if $|\mathbf{x}_1 \cdot \mathbf{u} - \mathbf{x}_2 \cdot \mathbf{u}| > r$, $h(\mathbf{x}_1) \neq h(\mathbf{x}_2)$.

Analysis of the Hashing Scheme (Cont.)

Define $Z = \mathbf{x}_1 \cdot \mathbf{u} - \mathbf{x}_2 \cdot \mathbf{u}$ and $d = \|\mathbf{x}_1 - \mathbf{x}_2\|_2$. Then, Z follows a Gaussian distribution of $\mathcal{N}(0, d^2)$.

Then, the probability distribution of the event $h(\mathbf{x}_1) = h(\mathbf{x}_2)$ is:

$$\mathbb{P}[h(\mathbf{x}_1) = h(\mathbf{x}_2)] = \int_0^r \mathbb{P}[h(\mathbf{x}_1) = h(\mathbf{x}_2) | |Z| = y] \cdot f(y) dy,$$

where $f(y)$ (with $y \geq 0$) is the density function of $|Z|$. Since $\mathbb{P}[h(\mathbf{x}_1) = h(\mathbf{x}_2) | |Z| = y] = 1 - \frac{y}{r}$, we have that:

$$\begin{aligned} \mathbb{P}[h(\mathbf{x}_1) = h(\mathbf{x}_2)] &= \int_0^r \left(1 - \frac{y}{r}\right) \cdot f(y) dy \\ &= \int_0^r \frac{1}{d} f\left(\frac{y}{d}\right) \left(1 - \frac{y}{r}\right) dy, \end{aligned}$$

where $f(y)$ is the density function of $|\mathcal{N}(0, 1)|$.

Analysis of the Hashing Scheme (Cont.)

Now we make a connection between the distance of $\mathbf{x}_1$ and $\mathbf{x}_2$ and the probability that they will collide. Define the probability derived on previous page as $p(\|\mathbf{x}_1 - \mathbf{x}_2\|)$.

- Clearly, $p_1 = p(r)$;
- $p_2 = p(c \cdot r)$.

We want to bound $\rho = \log(1/p(r))/\log(1/p(c \cdot r))$. Here, notice that for different data points $\mathbf{x}_1$ and $\mathbf{x}_2$, p_1 and p_2 clearly depend on these two inputs and can be different for different input $\mathbf{x}_1$ and $\mathbf{x}_2$.

Lemma 6 ([2])

Given the LSH designed with Equation 1, $\rho = \frac{\log(1/p_1)}{\log(1/p_2)}$ is bounded by $\frac{1}{c}$.

Table of Contents

1 Locality Sensitive Hashing

2 LSH Family for Different Distance Measures

- LSH for Euclidean distance
- LSH for Cosine Distance
- LSH for Jaccard Distance (for Binary Vectors)

Cosine Distance

Cosine similarity between two vectors $\mathbf{x}$ and $\mathbf{y}$:

$$\cos(\mathbf{x}, \mathbf{y}) = \frac{\mathbf{x} \cdot \mathbf{y}}{\|\mathbf{x}\|_2 \cdot \|\mathbf{y}\|_2}$$

The **cosine distance** measures the angle between the two vectors:

$$\text{dist}_{\cos}(\mathbf{x}, \mathbf{y}) = \arccos\left(\frac{\mathbf{x} \cdot \mathbf{y}}{\|\mathbf{x}\|_2 \cdot \|\mathbf{y}\|_2}\right).$$

Easy to verify:

- When $\mathbf{x}$ and $\mathbf{y}$ point to opposite directions, their cosine distance is the largest, which is π .
- When $\mathbf{x}$ and $\mathbf{y}$ point to the same direction, their cosine distance is the smallest, which is 0.

LSH for Cosine Distance: SimHash [3]

SimHash is designed for the cosine distance. Given a sampled random Gaussian unit vector $\mathbf{u}$, $h_{\mathbf{u}}(\mathbf{x})$ is defined as:

$$h_{\mathbf{u}}(\mathbf{x}) = \text{sign}(\mathbf{u} \cdot \mathbf{x}),$$

where sign outputs 1 if the input is ≥ 0 and otherwise outputs -1.

Lemma 7

Given two data points $\mathbf{x}$ and $\mathbf{y}$ and a SimHash function $h_{\mathbf{u}}(\cdot)$. The probability that the hash values of $\mathbf{x}$ and $\mathbf{y}$ are equal, i.e., either both are 1 or both are -1, is as follows:

$$\mathbb{P}[h_{\mathbf{u}}(\mathbf{x}) = h_{\mathbf{u}}(\mathbf{y})] = 1 - \frac{\text{dist}_{\text{cos}}(\mathbf{x}, \mathbf{y})}{\pi}.$$

Analysis of SimHash

We can prove Lemma 7 by drawing the figure as shown on the right-hand side.

We further consider

$$\rho = \log(1/p_1)/\log(1/p_2).$$

$$\rho = \frac{\ln(1/(1-r/\pi))}{\ln(1/(1-cr/\pi))}$$

By using the following inequality:

$$\frac{x}{x+1} \leq \ln(1+x) \leq \frac{x(6+x)}{6+4x} \leq x \text{ for } x > -1.$$

We can show that $\rho \leq \frac{1}{c}$ for $c \geq 2$.




Table of Contents

1 Locality Sensitive Hashing

2 LSH Family for Different Distance Measures

- LSH for Euclidean distance
- LSH for Cosine Distance
- LSH for Jaccard Distance (for Binary Vectors)

Jaccard Distance for Binary Vectors

Given two binary vectors $\mathbf{x}$ and $\mathbf{y}$, we first map them to two sets X and Y , where the i -th non-zero entry indicates the existence of element i in the set. The Jaccard similarity $J(\mathbf{x}, \mathbf{y})$ is defined as:

$$J(\mathbf{x}, \mathbf{y}) = \frac{|X \cap Y|}{|X \cup Y|}.$$

The Jaccard distance $\text{dist}_{\text{Jaccard}}(\mathbf{x}, \mathbf{y})$ is defined as $1 - J(\mathbf{x}, \mathbf{y})$.

Example 1

Let $\mathbf{x} = (1, 0, 0, 0, 1, 0, 1, 0, 1, 1)$ and $\mathbf{y} = (0, 0, 0, 0, 1, 0, 1, 0, 0, 0)$. Then, the Jaccard similarity between $\mathbf{x}$ and $\mathbf{y}$ is

$$J(\mathbf{x}, \mathbf{y}) = \frac{|X \cap Y|}{|X \cup Y|} = \frac{\{1, 5, 7, 9, 10\} \cap \{5, 7\}}{\{1, 5, 7, 9, 10\} \cup \{5, 7\}} = \frac{2}{5}$$

Then, $\text{dist}_{\text{Jaccard}}(\mathbf{x}, \mathbf{y}) = 1 - J(\mathbf{x}, \mathbf{y}) = \frac{3}{5}$.

MinHash for Jaccard Distance [4]

Generate a random permutation P for $[1 \cdots d]$. For the i -th dimension, the hashed value is $P[i]$, i.e., the value in the i -th dimension of the permutation P . Let $h(\mathbf{x}) = \min_{x(i)=1} P(i)$ be the minimum hashed value among the non-zero entries. Then,

$$\mathbb{P}[h(\mathbf{x}) = h(\mathbf{y})] = J(\mathbf{x}, \mathbf{y}) = 1 - \text{dist}_{\text{Jaccard}}(\mathbf{x}, \mathbf{y}).$$

Example 2

Let $\mathbf{x} = (1, 0, 0, 0, 1, 0, 1, 0, 1, 1)$ and $\mathbf{y} = (0, 0, 0, 0, 1, 0, 1, 0, 0, 0)$. Assume that the random permutation of $[1 \cdots 10]$ is $[4, 2, 10, 5, 1, 3, 8, 7, 9, 6]$. Then, $h(\mathbf{x}) = \min\{4, 1, 8, 9, 6\} = 1$ and $h(\mathbf{y}) = \min\{1, 8\} = 1$.

We can verify that $\rho = \log(1/p_1)/\log(1/p_2)$ can be bounded by $\frac{1}{c}$ if $c \geq 2$.

References



Piotr Indyk and Rajeev Motwani.

Approximate nearest neighbors: Towards removing the curse of dimensionality.
In *STOC*, pages 604–613, 1998.



Mayur Datar, Nicole Immorlica, Piotr Indyk, and Vahab S. Mirrokni.

Locality-sensitive hashing scheme based on p-stable distributions.
In *SOCG*, pages 253–262, 2004.



Moses Charikar.

Similarity estimation techniques from rounding algorithms.
In *STOC*, pages 380–388, 2002.



Andrei Z. Broder, Moses Charikar, Alan M. Frieze, and Michael Mitzenmacher.

Min-wise independent permutations (extended abstract).
In *STOC*, pages 327–336, 1998.